

HRUSHIKESH WADEKAR

Mumbai, India | +91 7249224098 | wadekarhrushikesh3@gmail.com | linkedin.com/in/hrushikeshwadekar

PROFESSIONAL SUMMARY

Certified Data Engineer (**Snowflake SnowPro Core, AWS Solutions Architect, Databricks Data Engineer**) with 2+ years of experience engineering high-throughput batch ETL/ELT pipelines in the **BFSI / Banking** domain. Proven expertise at **Tata Consultancy Services (Client: State Bank of India)** processing **4.45M+ daily transaction records** across 12-hour settlement cycles using **PySpark, SQL, Apache Airflow, and Oracle DB**. Skilled in resolving Spark Out-Of-Memory (OOM) bottlenecks, architecting Kimball dimensional models, and deploying automated financial reconciliation frameworks with 99.8% SLA reliability.

CERTIFICATIONS

- | | |
|---|------------------|
| • Snowflake SnowPro Core Certification — <i>Snowflake Inc.</i>
ID: S159845-260728-COF | Issued: Jul 2026 |
| • AWS Certified Solutions Architect – Associate — <i>Amazon Web Services (AWS)</i>
ID: fd461ec57a104b2381260114d1a42bf1 | Issued: Jun 2026 |
| • Databricks Certified Data Engineer Associate — <i>Databricks</i>
ID: 193994560 | Issued: Sep 2026 |

TECHNICAL SKILLS

Big Data & Processing: Apache Spark (PySpark, Spark SQL), Broadcast Joins, Partitioning, Shuffle Optimization, OOM Debugging
Workflow Orchestration: Apache Airflow (DAG Design, Scheduling, Task Dependencies, Failure Alerting, SLA Monitoring)
Cloud & Warehouses: Snowflake (Virtual Warehouses, Data Modeling, SnowSQL), AWS (AWS S3, Cloud Fundamentals), Databricks
Databases & Querying: Oracle Database (Complex Queries, Joins, CTEs, Window Functions, PL/SQL, Performance Tuning)
Data Modeling & Design: Dimensional Modeling (Kimball Star Schema, Snowflake Schema), Fact & Dimension Tables, SCD Type 1 & 2
Languages & Scripting: Python (ETL Scripting, Data Structures), SQL (Advanced Transformations), Bash / Shell Scripting
Data Quality & Governance: Pre/Post-load Financial Reconciliation, Data Validation Checks, Schema Consistency, Audit Compliance
Tools & Operating Systems: Linux

WORK EXPERIENCE

Data Engineer — <i>Tata Consultancy Services (Client: State Bank of India)</i> Project: Vendors Payment & Settlement System (VPS) <i>Mumbai, India</i>	Apr 2024 – Present
---	---------------------------

- **High-Throughput Batch Ingestion:** Architected and maintained end-to-end batch ETL pipelines ingesting **4.45M+ transaction hits** across 12-hour daily settlement cycles from core banking Oracle staging environments into analytics datamarts.
- **Airflow Orchestration & SLA Adherence:** Scheduled and monitored **20+ production Apache Airflow DAGs**, implementing robust task dependencies and automated alerting mechanisms that sustained a **99.8% workflow success rate**.
- **Spark Performance Tuning & RCA:** Diagnosed and eliminated recurring PySpark Out-of-Memory (OOM) exceptions and data skew by implementing **broadcast joins, optimal repartitioning, and shuffle partition tuning**, cutting pipeline downtime by **30%**.
- **Financial Ledger Reconciliation:** Engineered automated **pre-load and post-load reconciliation scripts** in SQL and PySpark to cross-verify debit/credit ledger balances and transaction counts, driving a **25% improvement in data accuracy** for audit compliance.
- **Database & Query Optimization:** Refactored complex SQL transformations, indexing, and batch loading procedures in Oracle DB, slashing batch execution runtimes by **20%** and improving reporting throughput.

TECHNICAL PROJECTS

Automated REST API Batch Ingestion & Dimensional Warehouse Pipeline

Technologies: Python, PySpark, Snowflake, AWS S3, SQL, Dimensional Modeling (Star Schema)

- Built a modular batch ingestion framework in **Python** fetching semi-structured JSON payloads from multi-endpoint REST APIs with automated pagination, token handling, and robust retry mechanisms.
- Executed scalable distributed transformations using **PySpark** to clean, flatten nested JSON schemas, and enforce schema validation rules prior to staging on **AWS S3**.
- Designed and deployed a Kimball **Star Schema** data warehouse in **Snowflake**, structuring normalized dimension tables (SCD Type 1) and transactional fact tables for business intelligence queries.
- Implemented automated data quality assertions verifying row counts, primary key uniqueness, and null constraints across analytical datasets.

EDUCATION

Master of Computer Applications (MCA) G. H. Raisoni College of Engineering and Management (KBC NMU), Jalgaon	2025 – 2027 (Pursuing)
Bachelor of Computer Applications (BCA) R. C. Patel Arts, Commerce and Science College (KBC NMU), Shirpur CGPA: 9.62 / 10.0	2020 – 2023