

HRUSHIKESH WADEKAR

Mumbai, Maharashtra, India | +91 7249224098 | wadekarhrushikesh3@gmail.com | linkedin.com/in/hrushikeshwadekar

PROFESSIONAL SUMMARY

Databricks and Microsoft Fabric Certified Data Engineer and AWS Solutions Architect with 2+ years of production experience at **Tata Consultancy Services (Client: State Bank of India)** building high-throughput batch ETL/ELT pipelines, distributed data transformations, and analytical datamarts in the **BFSI / Banking** domain. Proven expertise processing **4.45M+ daily transaction records** across 12-hour settlement cycles using **PySpark, Spark SQL, Apache Airflow, and Oracle DB**. Skilled in diagnosing Spark Out-Of-Memory (OOM) bottlenecks, optimizing shuffle partitions and broadcast joins, and implementing automated reconciliation frameworks with 99.8% SLA reliability.

CERTIFICATIONS

- Snowflake SnowPro Core Certification** — *Snowflake Inc.* (Issued: Jul 2026)
ID: S159845-260728-COF
- AWS Certified Solutions Architect – Associate** — *Amazon Web Services (AWS)* (Issued: Jun 2026)
ID: fd461ec57a104b2381260114d1a42bf1
- Databricks Certified Data Engineer Associate** — *Databricks* (Issued: Sep 2026)
ID: 193994560
- Microsoft Certified: Fabric Data Engineer Associate** — *Microsoft* (Issued: Sep 2026)
ID: F2CEA18F28A701DC

TECHNICAL SKILLS

Big Data & Processing: Apache Spark (PySpark, Spark SQL), Broadcast Joins, Partitioning, Shuffle Optimization, OOM Debugging
Cloud & Warehouses: Databricks, Microsoft Fabric (OneLake), AWS (AWS S3, Cloud Fundamentals), Snowflake (Virtual Warehouses, SnowSQL)
Workflow Orchestration: Apache Airflow (DAG Design, Scheduling, Task Dependencies, Failure Alerting, SLA Monitoring)
Languages & Scripting: Python (ETL Scripting, Data Structures), SQL (Advanced Transformations, Complex Queries), Bash / Shell Scripting
Databases & Querying: Oracle Database (PL/SQL, Joins, CTEs, Window Functions, Performance Tuning)
Data Quality & Governance: Pre/Post-load Financial Reconciliation, Data Validation Checks, Schema Consistency, Audit Compliance
Data Modeling & Design: Dimensional Modeling (Kimball Star Schema, Snowflake Schema), Fact & Dimension Tables, SCD Type 1 & 2
Tools & Operating Systems: Linux, Git

WORK EXPERIENCE

Data Engineer — *Tata Consultancy Services (Client: State Bank of India)* **Apr 2024 – Present**
Mumbai, India
Project: Vendors Payment & Settlement System (VPS)

- Spark Performance Tuning & RCA:** Diagnosed and eliminated recurring PySpark Out-of-Memory (OOM) exceptions and data skew by implementing broadcast joins, optimal repartitioning, and shuffle partition tuning, cutting pipeline downtime by 30%.
- High-Throughput Batch Ingestion:** Architected and maintained end-to-end batch ETL pipelines ingesting 4.45M+ transaction hits across 12-hour daily settlement cycles from core banking Oracle staging environments into analytics datamarts.
- Airflow Orchestration & SLA Adherence:** Scheduled and monitored 20+ production Apache Airflow DAGs, implementing robust task dependencies and automated alerting mechanisms that sustained a 99.8% workflow success rate.
- Database & Query Optimization:** Refactored complex SQL transformations, indexing, and batch loading procedures in Oracle DB, slashing batch execution runtimes by 20% and improving reporting throughput.
- Financial Ledger Reconciliation:** Engineered automated pre-load and post-load reconciliation scripts in SQL and PySpark to cross-verify debit/credit ledger balances and transaction counts, driving a 25% improvement in data accuracy for audit compliance.

TECHNICAL PROJECTS

Automated REST API Batch Ingestion & Dimensional Warehouse Pipeline

Technologies: Python, PySpark, Snowflake, AWS S3, SQL, Dimensional Modeling (Star Schema)

- Built a modular batch ingestion framework in Python fetching semi-structured JSON payloads from multi-endpoint REST APIs with automated pagination, token handling, and robust retry mechanisms.
- Executed scalable distributed transformations using PySpark to clean, flatten nested JSON schemas, and enforce schema validation rules prior to staging on AWS S3.
- Designed and deployed a Kimball Star Schema data warehouse in Snowflake, structuring normalized dimension tables (SCD Type 1) and transactional fact tables for business intelligence queries.
- Implemented automated data quality assertions verifying row counts, primary key uniqueness, and null constraints across analytical datasets.

EDUCATION

Master of Computer Applications (MCA) **2025 – 2027 (Pursuing)**
G. H. Rasoni College of Engineering and Management (KBC NMU), Jalgaon

Bachelor of Computer Applications (BCA) **2020 – 2023**
R. C. Patel Arts, Commerce and Science College (KBC NMU), Shirpur | CGPA: 9.62 / 10.0